

# Autonomic Computer Systems CS321: Self-Stabilization, logical clocks, distr. snapshot, Byzantine agreement

September 27, 2011

Prof. Dr. Christian Tschudin

Departement Mathematik und Informatik, Universität Basel

## Overview


---

### Selected “Self-” Algorithms (I)

- From last session:
  - Notes on message passing models
  - intro self-stabilization
- **Self-stabilization (example: mutual exclusion)**
- Logical clocks
- Distributed snapshot
- Byzantine agreement

# Mutual Exclusion in a Ring

Also known as **Token Ring**:



(Thijs Krol, UTwente)

- Anonymous ring (no IDs)
- One special node '0', otherwise identical algorithm
- One node is owner ("has token", "privileged" node)
- Token is passed on clockwise
- Note: no need for all nodes to know who has the token (LE)

## Token Ring – Algorithm

$N$  = number of nodes, number  $m$  with:


each node  $i$  has state  $s(i) \in \{0, \dots, m-1\}$ ,  $m \geq N-1$



Node 0: is privileged if  $s(0) == s(N-1)$   
when releasing token do:  $s(0) := (s(0) + 1) \bmod m$

Node  $i$ : is privileged if  $s(i) != s(i-1)$   
when releasing token do:  $s(i) := s(i-1)$

## TR – Algorithm (contd 1: normal)



- Example shows initial consistent states: one token circulates.

- Algorithm:

Node 0: wins if  $s(0) == s(N-1)$

Release token:

$s(0) := (s(0) + 1) \bmod m$


Node i: wins if  $s(i) != s(i-1)$

Release token:

$s(i) := s(i-1)$

(Thijs Krol, UTwente)

## TR – Algorithm (contd 2: convergence)



- Example shows initial arbitrary states, but only one token survives.

- Algorithm:

Node 0: wins if  $s(0) == s(N-1)$

Release token:

$s(0) := (s(0) + 1) \bmod m$

Node i: wins if  $s(i) != s(i-1)$

Release token:

$s(i) := s(i-1)$

# Token Ring – Properties

---

- The algo above is **Self-Stabilizing**
  - arbitrary initial values for  $s(i)$
  - yet converges to correct 1-token situation
- Errors can be added, at any time
  - system will regain legal state
- Theorem: There is no Token Ring algorithm for an anonymous ring without special node 0.

## Note on Self-Stabilization

---

Problem (and first solution) was proposed by Dijkstra in 1973

- Systems with a selfstab algorithm can recover from “loss of coordination”  
(due to node crash, memory crash etc)
- SelfStab **does not mask failures** (e.g. reliable transport, fault tolerance), but **recovers** from failure.

Transition to next topic:

one LE algo requires “timestamps”. How to obtain?

## Selected “Self-” Algorithms (II)

- Notes on message passing models
- Self-stabilization (example: mutual exclusion)
- **Logical clocks**
- Distributed snapshot
- Byzantine agreement

## Logical clocks


---

No global “now”: all message exchange takes some finite time.

- Logical ordering of events:
  - “happens-before” relation
  - obvious inside a node
  - two nodes: send happens before receive
- Each node keeps local clock:
  - counts events (sending, receiving)
  - we also say: messages receive a timestamp at reception

Leads to partial order. Some events  $a, b$  are “concurrent” i.e., neither  $\text{time}(a) < \text{time}(b)$  nor  $\text{time}(b) > \text{time}(a)$


# Cause and Effect, Independent or Concurrent



## Logical Clocks

$L(\text{event})$ : assign a number to an event

- We want: if  $a \rightarrow b$ , then  $L(a) < L(b)$
- Logical clocks do this, but  $L(a) < L(b)$  does not imply  $a \rightarrow b$  (integers have total order, but happens-before is partial)



Problem when comparing g and b

# Vector Clocks (Mattern, 1989)

---

- Each node keeps a vector of all logical clocks it learned, vector is timestamp
- On reception, use max of component values
- $<$  defined such that:  
All components equal or less, and then at least one component must be different.
- Incomparable values! They mean “concurrent”.  
See “causality” figure.

## Overview

---

### Selected “Self-” Algorithms (II)

- Notes on message passing models
- Self-stabilization (example: mutual exclusion)
- Logical clocks
- **Distributed snapshot**
- Byzantine agreement

# Global State

No global time available (you can't reconstruct *exact* physical time over asynchronous channels), no accurate global state obtainable.


- Weaker form: global state = combination of local states of processes and channels **at some time which could have occurred**
- Why we want global state:
  - finding lost token
  - termination of distributed computation
  - garbage collection.

We assume a strongly connected network.

## “Consistent State”

Intuitively: **consistent global state** = “snapshot” that looks to all processes as if it was taken at the same instant, everywhere.

- Causal relations must be observed, **consistent cut**





# Consistent Cut

---

Global state = must be a consistent cut

Property:

- Either:  
all events before the cut have to “happened-before” the events after the cut,
- Or: the events are unrelated (“concurrent”)
- This can be checked with vector clocks.

The big problem are the messages in the channel, they are part of the state.

## Application of cuts

---

Nodes can crash – how to restart in a safe state?

- Logging:  
Nodes periodically write their current state to disk
- In case of a node crash, after restart, the node knows its latest “cut”
- Other nodes in the system have to roll back to a maximally consistent cut.

Other use of cuts: take a (distributed) snapshot! Next slides ...

# Distributed Snapshot

---

Snapshot algorithms to record a consistent state

- Snapshots can be used to detect stable states.  
Examples of stable states: lost token, deadlock, termination.
- One well-know algo: Chandy and Lamport  
It creates one (possible) consistent cut.
- Assumptions for Chandy and Lamport:
  - reliable message exchange, FIFO, unidirectional
  - no failures (of nodes, links)
  - strongly connected graph


## Chandy and Lamport

---

Idea: “drain” all messages pending in the channels,  
put a marker before and after these messages.

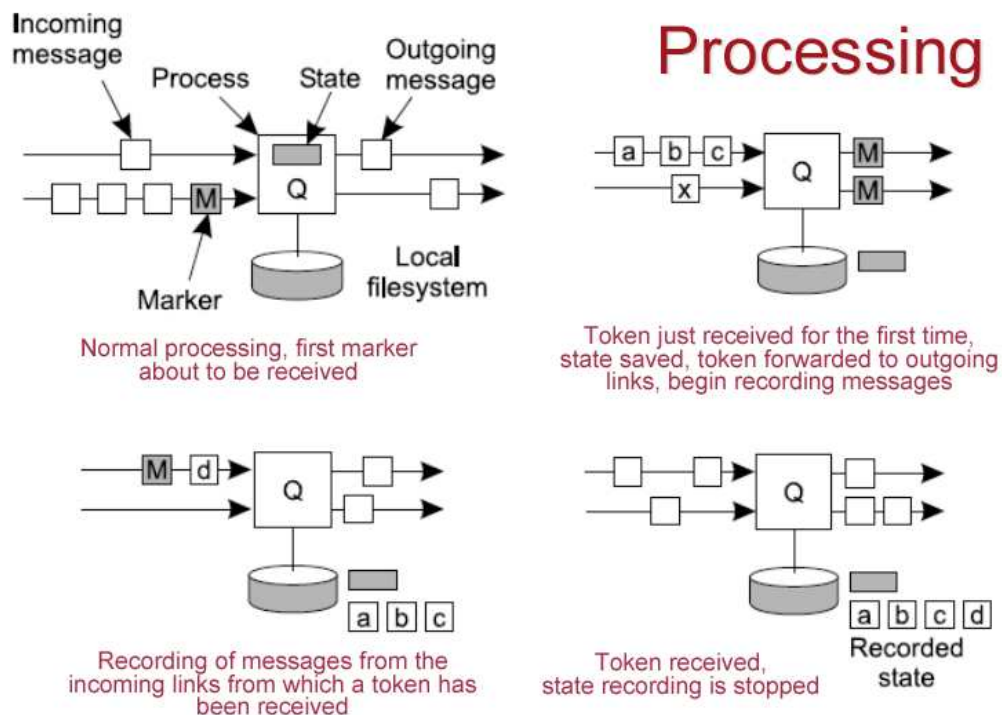
- Initiate snapshot by sending a marker M
- When receiving M
  - stop processing, record own state
  - send out marker on all links
  - record all subsequently received messages  
(except where M arrived)
  - until a next marker is received on each link.

# Chandy and Lamport: Intuition



“Token wave” which flushes the channels (two frontiers)

## Chandy and Lamport: Example at one Node



# Overview

---

Context: consensus finding (=distributed agreement)

- Notes on message passing models
- Self-stabilization (example: mutual exclusion)
- Logical clocks
- Distributed snapshot
- **Byzantine agreement**
  - two general's problem
  - byzantine failures
  - byzantine problem, (impossibility) results
  - the oral message protocol

## Byzantine Agreement and related topics

---

Presentation approach:

- Failure modes (including “byzantine failures”)
- Problematic consensus in case of message loss (“two generals problem”)
- Problematic consensus in a group, with reliable messages (byzantine agreement problem,  $3f + 1$  requirement)
- Impossibility result for asynchronous settings (Fischer, Lynch and Peterson, 1982)
- Discussion: ways out of this problem

# Failure Modes

---

Depending on the assumptions of failure, different algorithms (or impossibility results) will ensue.

- **Fail-stop:**
  - either a processor works correctly and participates
  - or it has failed and will never respond again.Moreover: the others can reliably detect failed processors.
- **Slowdown:** when some processors execute slowly (or fail), the others cannot know for sure.
- **Byzantine:** every action might be corrupt (see next slide)

Fail-stop is often assumed, but is not realistic (e.g. msg delay)

## The Byzantine Failure Mode

---

Byzantine failures include everything:

- lost messages
- crashing nodes
- faulty implementation
- **malicious implementation, even collusion!**

In the following we make things simple:

- synchronous execution model
- reliable message exchange
- full connectivity and sender authentication

That is: “only” the processors behave byzantine, it’s problem enough

## Prelude: “Two Generals’ Problem”

---

What if message passing is not reliable? Big trouble!

- Two generals have to coordinate an attack (“we attack at 4am”) – each must be absolutely sure that the other agrees – in case of doubt: none will attack
- Because of message loss, sender is uncertain – even with confirmation, this persists (the other thinks: did he receive my confirmation?)
- Result: impossible to get common agreement, for sure.

In practice: accept state with some probability  
(after N rounds we are “confident”, instead of sure)

## Byzantine Agreement Problem

---

The story:

- Romans still rule over Constantinople, but Ottoman battalions want to launch an attack.
- Only a coordinated attack will be successful, a coordinated retreat is also an option.
- The general decides to attack (or retreat): all lieutenants receive the order, lieutenants will exchange messages to verify the outcome.
- Some lieutenants (incl the general) are traitors, and pass on lies.

Is there a protocol for (successfully) coordination?

Also known as: **consensus algorithms, interactive consistency algorithms**

# Byzantine Agreement Problem, CS version

---

- Input: Each process starts with a bit
- Goal: run a protocol such that
  - all processes output the same bit
  - the output bit must match at least one of the initial bits
- Q: how many faulty processes can you tolerate?

The problem was proposed in 1980  
by Lamport, Shostak and Pease.

## Byzantine Agreement: Desired Properties

---

- Non-trivial protocol:  
if all lieutenants have same input  $b$ , the loyal lieutenants' agreement must be  $b$ .
- Core property:  
All loyal lieutenants agree on the same value  
(which must not be original  $b$ , but it must be the same for all)
- Protocol must terminate.

## Byzantine Agreement: Impossibility (1)


---

Theorem: **There is no protocol for byzantine agreement among 3 nodes if at least one node fails.**

- Applied to story:  
with three lieutenants, already one traitor lieutenant will prevent successful agreement
- Proof preparation:
  - we build a 6-node network of reliable processes
  - on this we **emulate** (malicious) failures
  - ie, some processes output (deliberate) wrong answers

## Byzantine Agreement: Impossibility (2)

---




(Figs from G. Candea)

Broadcast (left) is represented as point-to-point (right),  
the **blue processes  $P$**  run identical code as  $P$ .

Assumption:  $\exists$  byz. agr. protocol for three nodes and one failure



## Byzantine Agreement: Impossibility (3)




Given:  $P_2$  and  $P_3$  have initial input 0

What happens if  $P_1$  is faulty? E.g.  $P_1$  sends different values!

$P_2$  and  $P_3$  nevertheless agree on 0 (assumption!)

## Byzantine Agreement: Impossibility (4)




Given:  $P_1$  and  $P_2$  have initial input 1

What happens if  $P_3$  is faulty? E.g.  $P_3$  sends different values!

$P_1$  and  $P_2$  nevertheless agree on 1 (assumption!)

## Byzantine Agreement: Impossibility (5)




Given: **P1** and P3 have initial input 1 and 0, respect.  
What happens if P2 is faulty? E.g. P2 sends different values!

**P1** and P3 nevertheless agree on some value (assumption!)

## Byzantine Agreement: Impossibility (6)

Trying to harmonize the different views results in contradiction:



- **P1** and **P2** agreed on 1
- **P1** and P3 agreed on some value; because **P1** above went for 1, this must be 1 also for P3
- However: P3 already agreed on 0 (with P2)
- contradiction.

## Byzantine Agreement: Impossibility (6)

---

Another argument, graphical this time:




Fig. 1. Lieutenant 2 a traitor.




Fig. 2. The commander a traitor.

L1 does not know whom to obey: L1 cannot simply obey the commander because among the lieutenant(s), no consensus possible.

## Byzantine Agreement, Theory Results

---

- Theory result:  
Protocols exist, if we have  $3f+1$  nodes for  $f$  failures  
(see subsequent slides on OM)
- **But:** Result is different in asynchronous model:  
No protocols at all! A single faulty node is the end!  
The “celebrated result” of Fischer, Lynch and Peterson in 1982)

In 2007, more than 1200 follow-up publications ...

# Oral Message Protocol (1)

---

Proposed by Lamport et al.

- Idea of OM:  
repeat majority vote recursively in smaller and smaller groups
- (another idea would be:  
lieutenants not only exchange their bits, but also what they  
heard from whom)

Also called an interactive consistency protocol.

In the following:

Algorithm OM(0) for last round, OM(m) for previous rounds

# Oral Message Protocol (2)

---

OM(0) for terminating

- Commander sends his value to every lieutenant
- Each lieutenant uses the value received from commander

*Note: OM(k) means “this protocol tolerates up to k traitors”*

## Oral Message Protocol (3)


---

OM( $m$ ),  $m > 0$

- Commander sends his value to every lieutenant ( $v_i$ )
- Each lieutenant acts as a commander for OM( $m-1$ ): sends  $v_i$  to the other  $n - 2$  lieutenants
- Lieutenant  $i$  receives  $v_j$ , and for  $j \neq i$ , computes majority: this becomes his value for OM( $m$ )

### Example 1: One traitor lieutenant


---



OM(1): This time, a majority vote among the lieutenants works, as the loyal ones outnumber the malicious lieutenant.

## Example 2: The commander is a traitor

---



OM(1) works even if the commander is not loyal.

## Message complexity of the OM algorithm

---

Overall:  $O(n^m)$  invocations (!),  
where  $n$  is the number of processes,  $m$  number of faulty processes

- OM( $m$ ) invokes OM( $m-1$ )  $n - 1$  times
- each OM( $m-1$ ) invokes OM( $m-2$ )  $n - 2$  times
- etc

# Protocol variations (ways out of impossibility result)

---

Approach: change some assumptions, or the rules/tricks

randomized instead of deterministic algorithm, or private channels, or signed messages, or failure detectors (instead of proceeding despite failures).

- Ben-Or (1983): exponential (in time) asynchronous byzantine agreement, works for  $f < n/5$
- still open how much this can be improved