

# CS321 - Autonomic Computer Systems

## Introduction to Control Theory

2008-10-02

Christian Tschudin, University of Basel

Using slides from:

J. L. Hellerstein, IBM

Z. Gao, Cleveland State University

C. Lu, University of Virginia

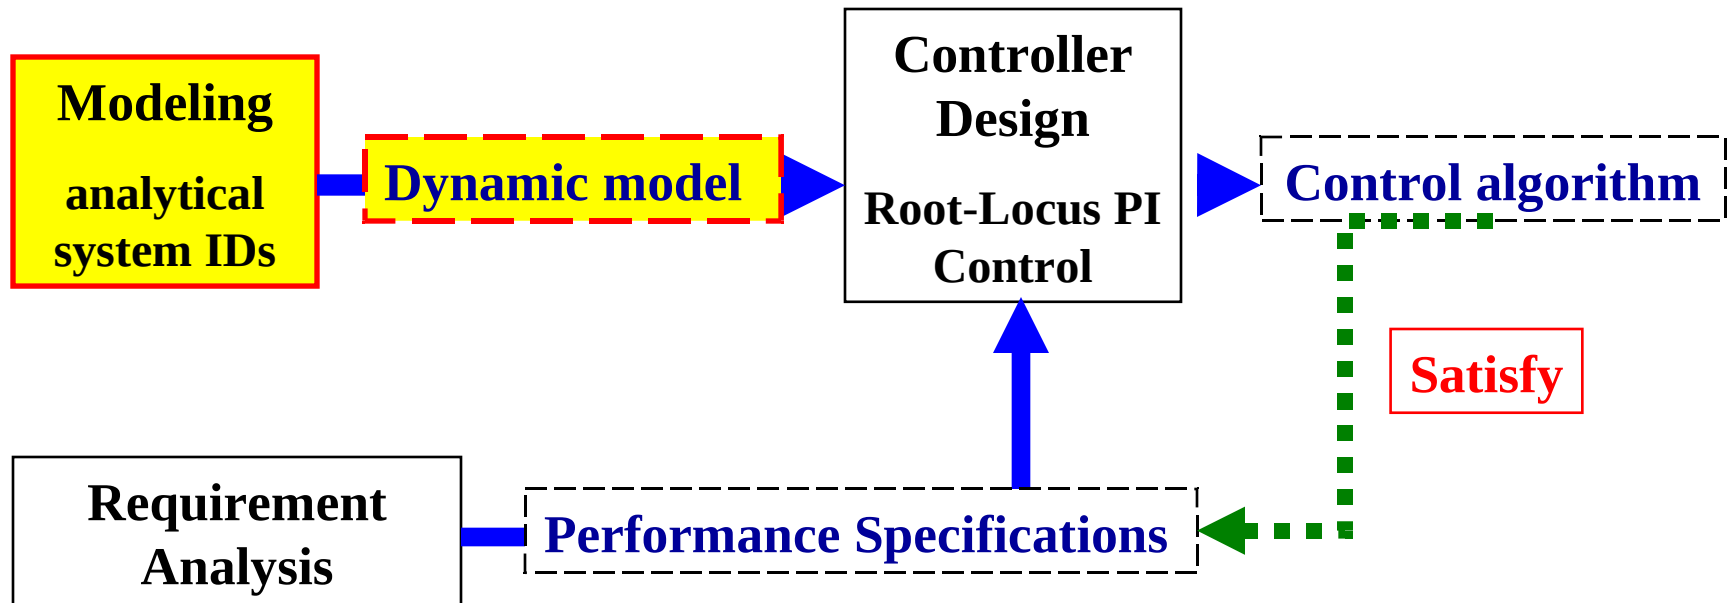
# Overview

- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# Feedback control theory, vs ...

- Adaptive resource management heuristics
  - ❖ Laborious design/tuning/testing iterations
  - ❖ Not enough confidence in face of untested workload
- Queuing theory
  - ❖ Doesn't handle feedbacks
  - ❖ Not good at characterizing transient behavior in overload
- Feedback control theory
  - ❖ Systematic theoretical approach for analysis and design
  - ❖ **Predict system response and stability to input**

# Control design methodology



# System Models

- **System ID** vs. First Principle

# Feedback Control of Computing Systems

## *M1: Introduction*

Joseph L. Hellerstein  
*IBM Thomas J Watson Research Center, NY*  
*[hellers@us.ibm.com](mailto:hellers@us.ibm.com)*

September 21, 2004

### Book reference:

“Feedback Control of Computing Systems”, Hellerstein, Diao, Parekh, Tilbury, Wiley, 2004.


Slides: [http://swig.stanford.edu/~fox/cs444a/ct\\_tutorial/](http://swig.stanford.edu/~fox/cs444a/ct_tutorial/)

# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

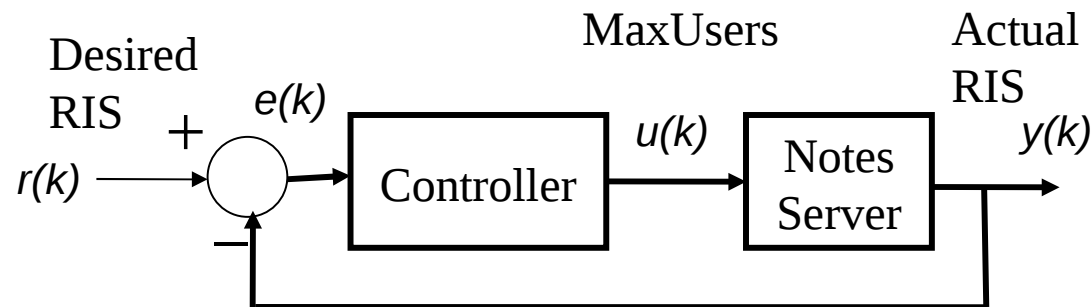
# Example: Control of Lotus Notes

## Architecture



RIS = RPCs in System

## Control Model



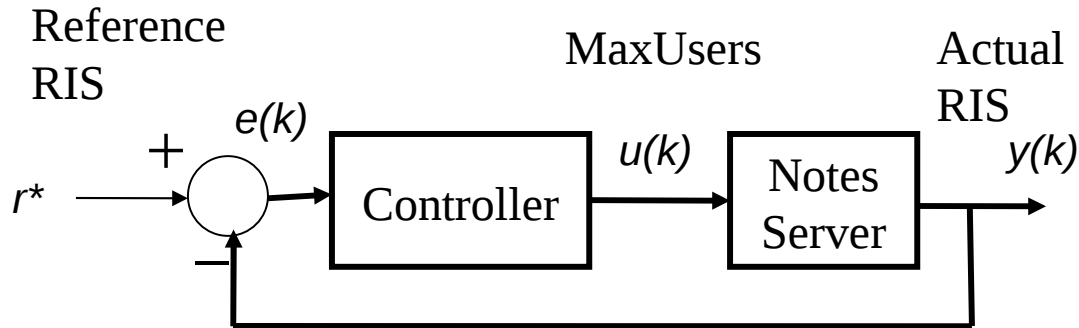
Control error:  $e(k) = r(k) - y(k)$

System model:  $y(k) = (0.43)y(k-1) + (0.47)u(k-1)$



# Lab Exercise

## Block Diagram



## ARX\* Models

Control error:  $e(k) = r^* - y(k)$

P controller:  $u(k) = Ke(k)$


System model:  $y(k) = (0.43)y(k-1) + (0.47)u(k-1)$

### ■ Spreadsheet file ctclass-M1


- ❖ **Proportional** controller:  $Maxusers(k+1) = K(r^* - y(k))$
- ❖ What is the effect of  $K$  on
  - Accuracy: (want  $r^* = y(k) = 200$ )
  - Stability
  - Convergence rate
  - Overshoot

\*ARX is autoregressive with an external input

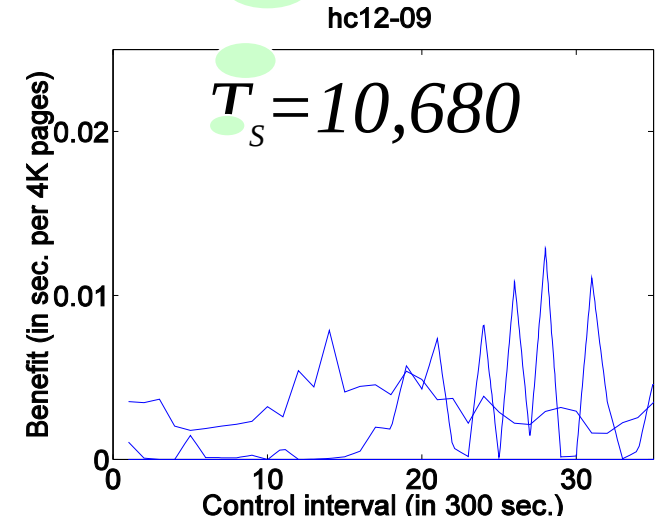
# Application of CT to a DBMS



**59%  
Reduction  
in Total RT**

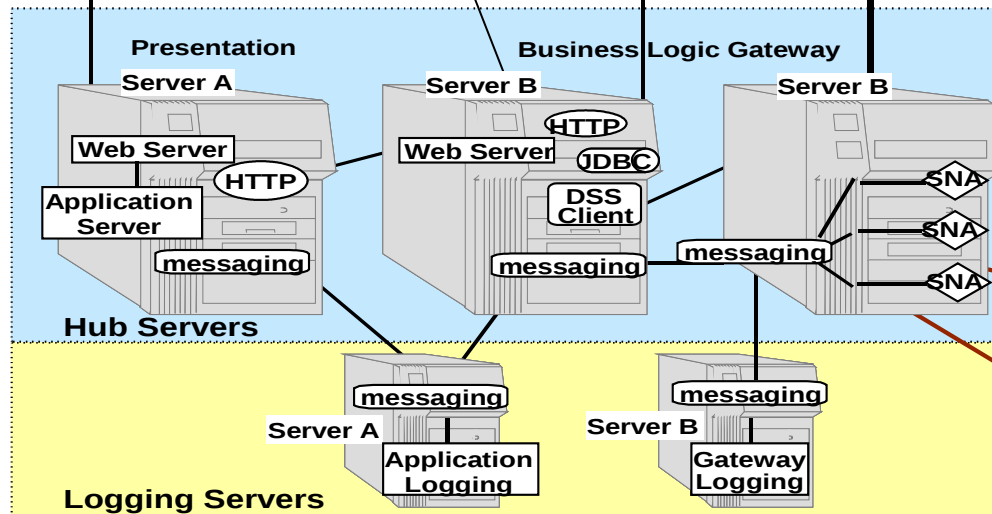



Without Controller




With Controller

# Why Control Theory?




Front end for online customer service

Many types of servers and applications




Back-end Systems

## Unstable System



- Stabililty
- Accuracy
- Settling time
- Overshoot

# Control System Architecture



## Components

Target system: what is controlled

Controller: exercises control

Transducer: translates measured outputs

## Data

Reference input: objective

Control input: manipulated to affect output

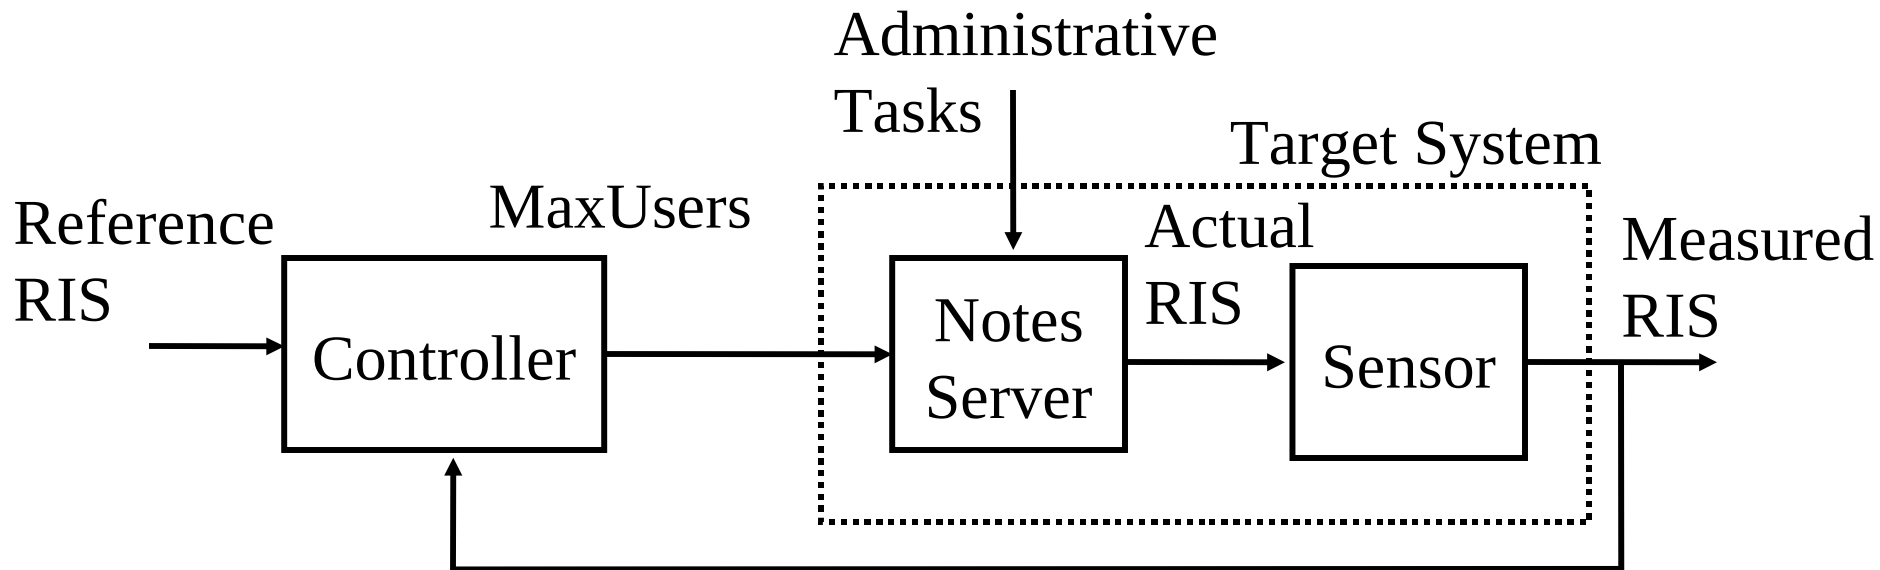
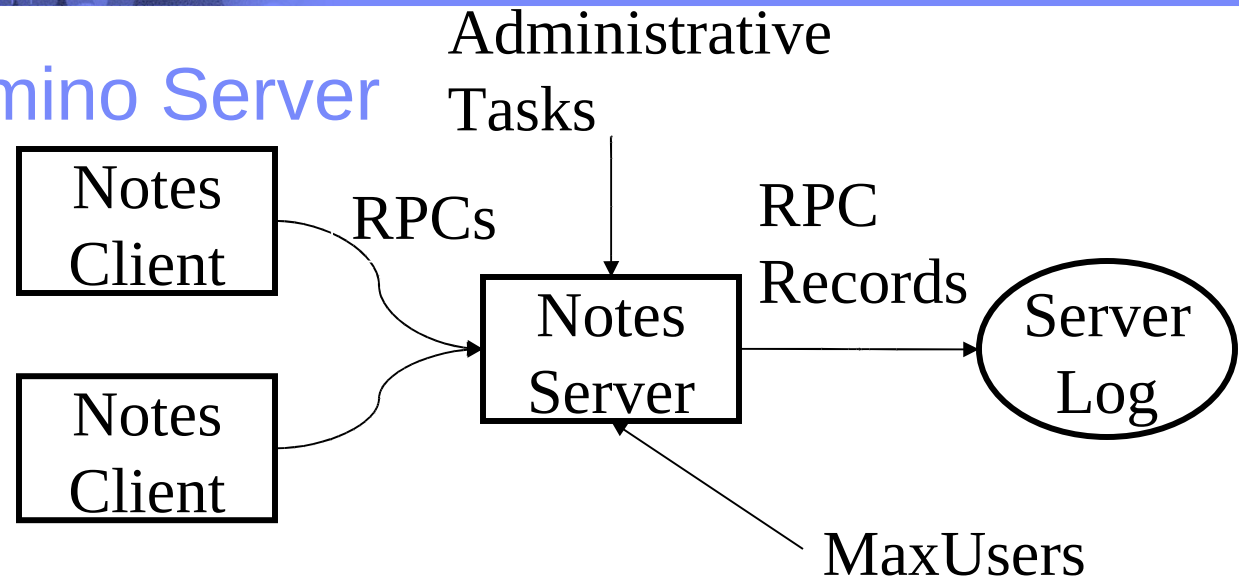
Disturbance input: other factors that affect the target system

Transduced output: result of manipulation

Given target system, transducer  
Control theory finds controller  
that adjusts control input  
to achieve measured  
output in the presence of  
disturbances.

# IBM Lotus Domino Server

Architecture



Block Diagram

# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# Properties of Control Systems



Stability

Accuracy


Short settling

Small overshoot

Unstable System

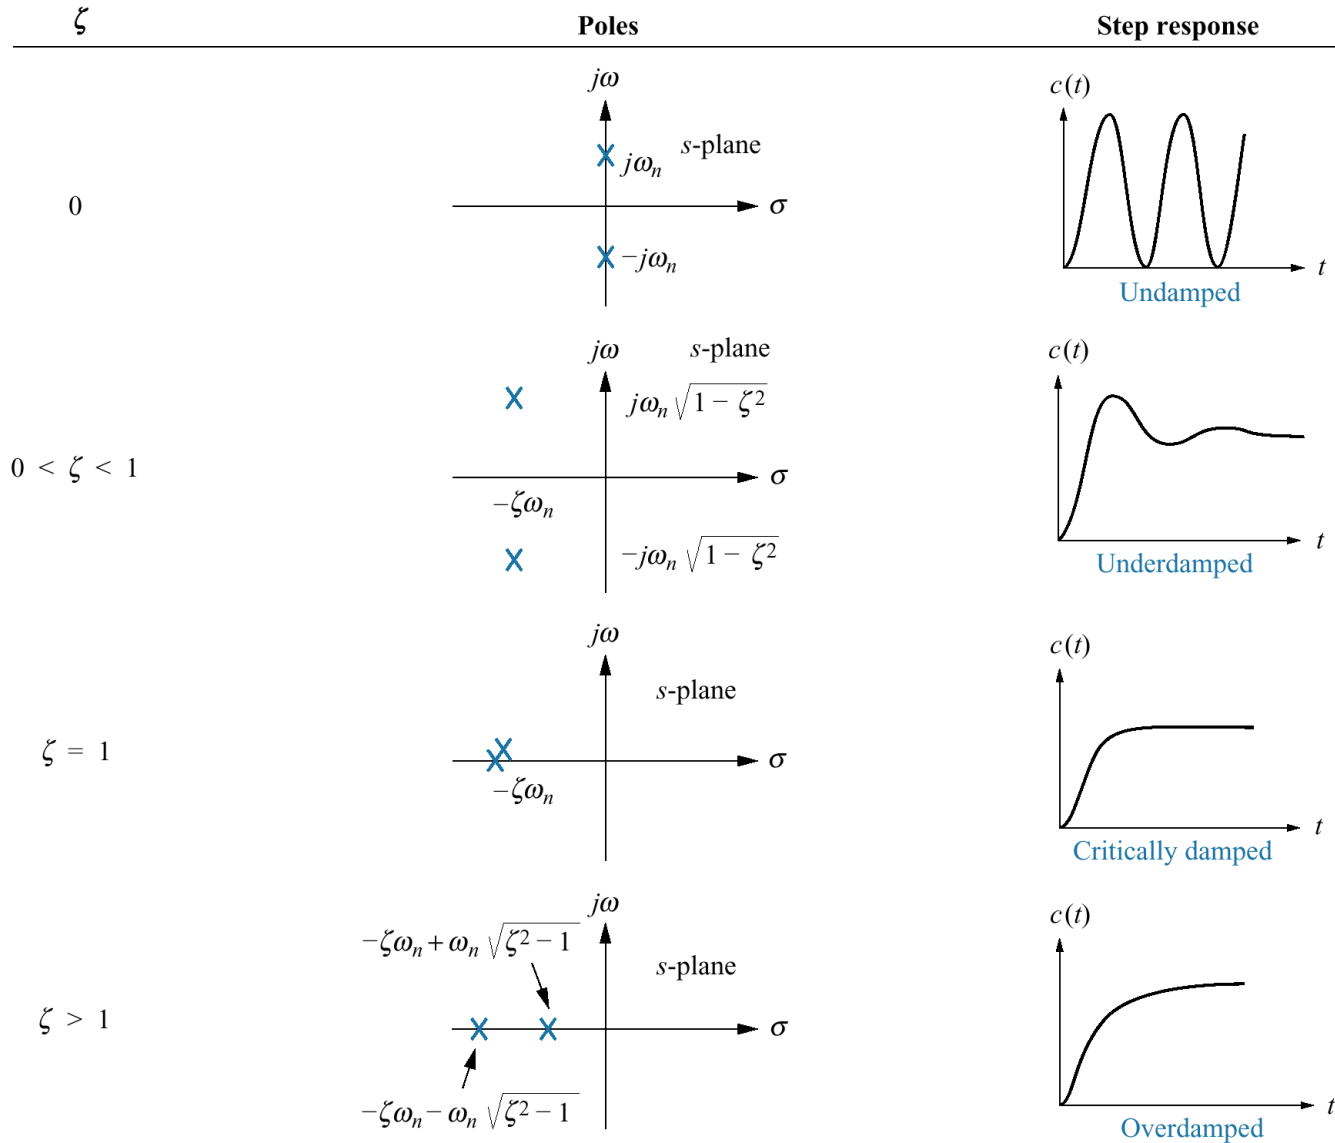


# Performance specifications





# Different “Responses” (Pure 2<sup>nd</sup> Order Transfer Functions)




# Control Theory in Two Slides: System Identification



Model of System Dynamics


$$y(k) = a_1 y(k-1) + b_1 u(k-1)$$



Transfer Function

$$N(z) = \frac{b_1}{z - a_1}$$

# Control Theory in Two Slides: Control Design







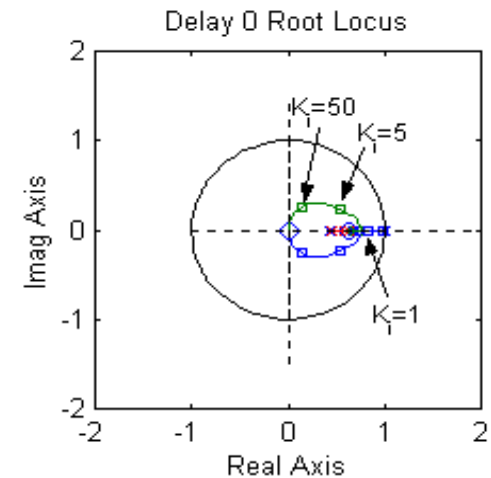


$H(z)$  = Closed Loop Transfer Function









Integral Control Law

$$u(t) = u(t-1) + Ke(t)$$

Poles  
of  
 $H(z)$



# Example: Control & Response in an Email Server



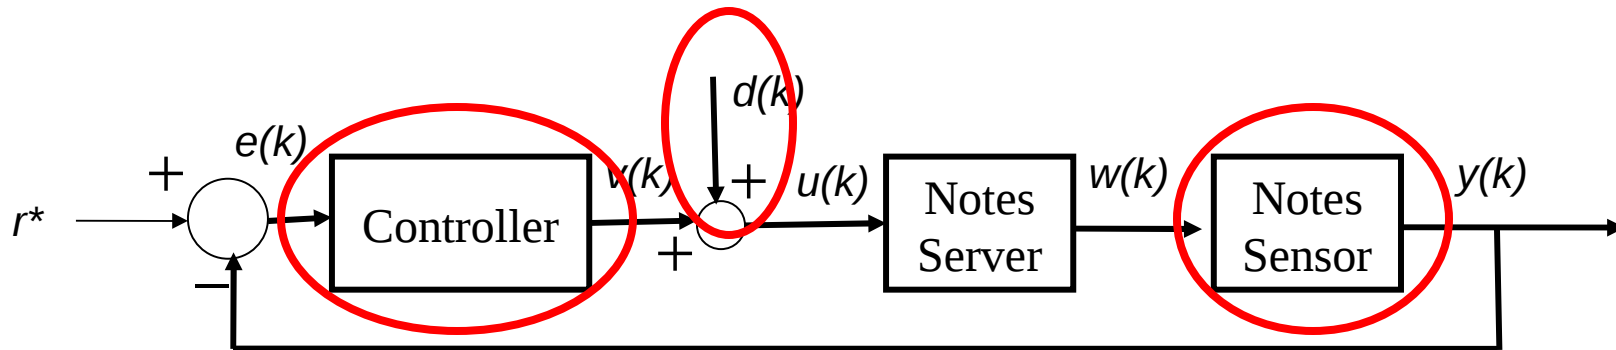
Good

Bad

Slow

Useless

## Lab 2: Modifications to the Lotus Notes Simulation



$$K=1, r(k)=200, d(k)=10$$


Control error:  $e(k) = r^* - y(k)$

I controller:  $v(k) = v(k-1) + Ke(k)$

System model:  $w(k) = (0.43)w(k-1) + (0.47)u(k-1)$

Disturbance input:  $u(k) = v(k) + d(k)$

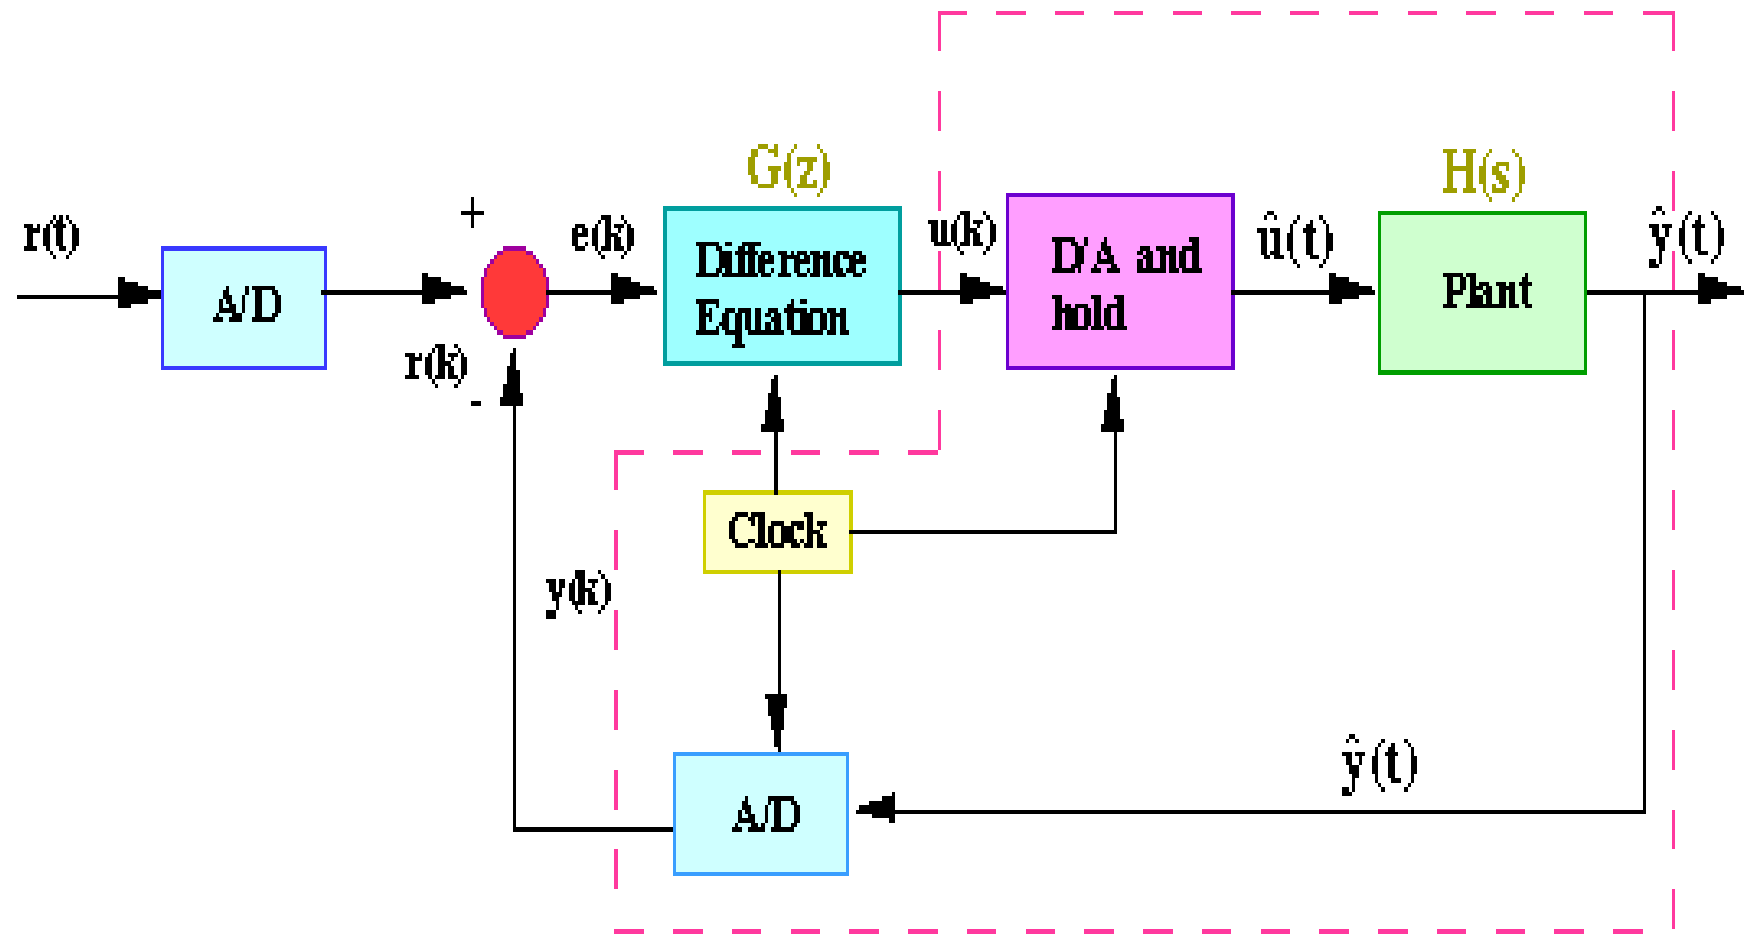
Sensor:  $y(k) = (0.8)y(k-1) + 0.72w(k) - 0.66w(k-1)$



# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# Digital Control



# Digital controler


- Discrete time
- Discrete values
- Usually: continuous time, continuous values
  - differential equations
  - use of Laplace transform
- With digital controler:
  - difference equations
  - use of z-Transform



# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- **z-Transform:**
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# Z-Transform of a Signal



## Time domain representation

$$u(0)=1$$

$$u(1)=3$$

$$u(2)=2$$

$$u(3)=5$$

$$u(4)=6$$

## z domain representation

$$1z^0 +$$

$$3z^{-1} +$$

$$2z^{-2} +$$

$$5z^{-3} +$$

$$6z^{-4}$$

$z^0 = 1$ :  $k = 0$  (current time)

$z^{-1}$ :  $k = 1$  (one time unit in the future)


$z^{-2}$ :  $k = 2$  (two time units in the future)

$z$  is time shift;  $z^{-1}$  is time delay  $\Rightarrow$


If  $\{u(k)\} = u(0), u(1), \dots$  is a signal, then its z-Transform is

$$U(z) = \sum_{k=0}^{\infty} u(k)z^{-k}$$

# Write the Z-Transforms or Plots for the Following Finite Signals




$$U(z) = -2 + 4z^{-1} + 5z^{-2} - z^{-3} + 4z^{-4}$$



$$V(z) = zU(z) = 4 + 5z^{-1} - z^{-2} + z^{-3}$$

(Drop exponents  $>0$ .)

## Common Signals: Impulse




$$y(0) = 1; y(k) = 0, k > 0$$

$$Y(z) = 1z^0 + 0z^{-1} + 0z^{-2} + \dots$$

$$= 1$$

$$\text{Example: } 5z^{-3} \Leftrightarrow u(3) = 5$$

## Common Signals: Unit Step $W(z)$



$$y(k) = 1, k \geq 0$$

$$Y(z) = 1z^0 + 1z^{-1} + 1z^{-2} + \dots$$

$$= \frac{1}{1 - z^{-1}}$$

$$= \frac{z}{z - 1}$$

# Properties of z-Transforms of Signals

Signals:  $U(z) = u(0)z^0 + u(1)z^{-1} + u(2)z^{-2} + \dots$

$$V(z) = v(0)z^0 + v(1)z^{-1} + v(2)z^{-2} + \dots$$

Shift:  $zU(z) = u(0)z^1 + u(1)z^0 + u(2)z^{-1} + \dots$

$$= u(1)z^0 + u(2)z^{-1} + \dots$$

Delay:  $U(z)/z = u(0)z^{-1} + u(1)z^{-2} + u(2)z^{-3} + \dots$

Scaling:  $aU(z) = au(0)z^0 + au(1)z^{-1} + au(2)z^{-2} + \dots$


$$= z\text{-Transform of } \{au(k)\}$$

Sum of signals:  $u(0)z^0 + u(1)z^{-1} + u(2)z^{-2} + \dots + v(0)z^0 + v(1)z^{-1} + v(2)z^{-2} + \dots$

$$= (u(0) + v(0))z^0 + (u(1) + v(1))z^{-1} + (u(2) + v(2))z^{-2} + \dots$$


$$= U(z) + V(z)$$

# Infinite Length Signals




$$\begin{aligned} U(z) &= 5 + 5z^{-1} + 5z^{-2} + \dots \\ &= \frac{5z}{z-1} \end{aligned}$$


# Construct a Z-Transform for the Following Infinite Length Signal




$Y(z)$



+



+



$$Y(z) = 5W(z) + 3z^{-10} + 5z^{-11}W(z)$$




# Common Signals: Geometric

Geometric:  $y(k) = a^k$

$$Y(z) = 1 + az + a^2 z^2 + \dots$$

$$= \frac{z}{z - a}$$



$$Y(z) = 1 + 0.8z^{-1} + 0.64z^{-2} + \dots$$

$$= \frac{z}{z - 0.8}$$

# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - **poles and stability**
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# Poles of a Z-Transform

Definition: Values of  $z$  for which the denominator is 0

Easy to find the poles of a geometric:  $V(z) = \frac{z}{z - a}$  Pole is  $a$ .

Quick exercises: What are the poles of the following Z-Transforms?

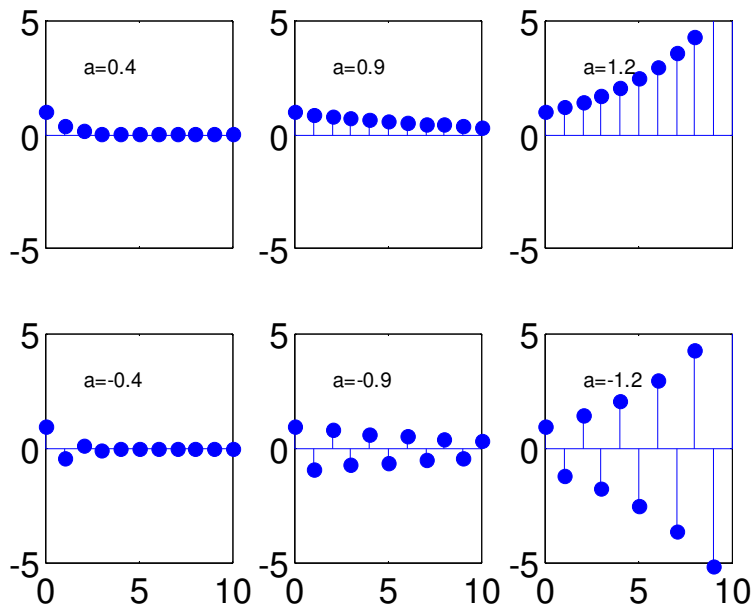
$Y(z) = \frac{5z}{z-1} + \frac{2z}{z-0.8}$  Easy if sum of geometrics

$V(z) = \frac{3z}{z^2 - 0.5z + 0.06}$  Harder if expanded polynomial

Poles determine key behaviors of signals

# Effect of Pole on the Signal

$$y(k) = a^k \Leftrightarrow \frac{z}{z-a}$$



Why?

$$\frac{z}{z-a} = 1 + az + a^2 z^2 + \dots \Leftrightarrow (1, a, a^2, \dots)$$

■ What happens when

■  $|a|$  is larger?

■  $|a| > 1$ ?

■  $a < 0$ ?

■ Larger  $|a|$

■ Slower convergence

■  $|a| > 1$

■ Does not converge

■  $a < 0$

■ Oscillates

# Final Value Theorem

- Provides an easy way to determine the steady state value of a signal

- ❖ Limit as  $k$  becomes large

$$v(\infty) = \lim_{z \rightarrow 1} (z-1)V(z)$$

(if  $V(z)$  has all of its poles inside the unit circle)

$$\lim_{z \rightarrow 1} (z-1) \frac{z}{z-1} = 1 \quad \text{Final value of the unit step is 1.}$$

$$\lim_{z \rightarrow 1} (z-1)1 = 0 \quad \text{Final value of the impulse is 0.}$$

$$\lim_{z \rightarrow 1} (z-1) \frac{z}{z-a} = 0, \text{ if } a \text{ is inside the unit circle}$$

# Overview



- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - **transfer function and system composition**
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# “Transfer Function” - Motivation and Definition

## Motivation:

ARX model relates  $u(k)$  to  $y(k)$

Transfer function expresses this relationship in the  $z$  domain



$$G(z) = \frac{\text{Output}}{\text{Input}} = \frac{Y(z)}{U(z)}$$



$$\text{or } Y(z) = G(z)U(z)$$

assuming initial conditions are 0.

**A transfer function is specified in terms of its input and output.**



Intuition for  $G(z)$ : Effect of “kicking the system”

# Constant Transfer Function





$$Y(z) = aU(z)$$

$$G(z) = \frac{Y(z)}{U(z)} = a$$







# Time-Delay Transfer Function





$$Y(z) = z^{-1}U(z)$$

$$G(z) = \frac{Y(z)}{U(z)} = z^{-1}$$





# Combining Simple Transfer Functions





$$Y(z) = az^{-1}U(z)$$

$$G(z) = \frac{Y(z)}{U(z)} = az^{-1}$$








# Additional Terms in T.F.










$$Y(z) = (az^{-1} + bz^{-2})U(z)$$

$$G(z) = \frac{Y(z)}{U(z)} = az^{-1} + bz^{-2}$$




# Geometric Sum of T.F.




$$Y(z) = (1 + az^{-1} + a^2z^{-2} + \dots)U(z)$$

$$G(z) = \frac{Y(z)}{U(z)} = 1 + az^{-1} + a^2z^{-2} + \dots = \frac{z}{z-a}$$





# T.F. Interpretation

Signal generated by an impulse input



Example:




$$G(z) = \frac{Y(z)}{U(z)}$$

$$Y(z) = G(z)U(z) = G(z)(1)$$

$$G(z) = \frac{z}{z-0.5} = 1 + 0.5z^{-1} + 0.25z^{-2} + 0.125z^{-3} + 0.0625z^{-4} + \dots$$

# Complicated Transfer Functions

## Decompose into a sum of geometrics



$$\begin{aligned} G(z) &= \frac{0.1z}{z^2 - 0.5z + 0.06} \\ &= \frac{z}{z - 0.3} - \frac{z}{z - 0.2} \\ &= 1 + 0.3z^{-1} + 0.09z^{-2} + \dots - 1 - 0.2z^{-1} - 0.04z^{-2} + \dots \\ &= 0.1z^{-1} + 0.05z^{-2} + \dots \end{aligned}$$

- **Partial fraction expansion** allows rational polynomials to be decomposed into a sum of geometrics
- Poles of the original polynomial are the poles of the geometrics

# Converting ARX Models into T.F.

## Approach

- Replace  $u(k+n)$  by  $z^n U(z)$  and replace  $y(k+n)$  by  $z^n Y(z)$
- Express  $Y(z)$  as a function of  $U(z)$
- Divide both sides by  $U(z)$

Example:

$$\text{T.S. Model: } y(k+1) = 0.43y(k) + 0.47u(k)$$

$$1: zY(z) = 0.43Y(z) + 0.47U(z)$$

$$2: (z - 0.43)Y(z) = 0.47U(z); Y(z) = \frac{0.47}{z - 0.43}U(z)$$

$$3: \frac{Y(z)}{U(z)} = \frac{0.47}{z - 0.43}$$

# Converting T.F. into ARX Model

## Approach

- Express transfer function as  $Y(z)=G(z)U(z)$
- Multiply both sides by the denominator of  $G(z)$
- Convert  $Y(z)$  and  $U(z)$  into the time domain

Example 1:

$$G(z) = \frac{z}{z - 0.3}$$

$$Y(z) = \frac{z}{z - 0.3} U(z)$$

$$Y(z)(z - 0.3) = zU(z)$$

$$zY(z) \Leftrightarrow y(k+1)$$

$$0.3Y(z) \Leftrightarrow 0.3y(k)$$

$$zU(z) \Leftrightarrow u(k+1)$$

$$y(k+1) - 0.3y(k) = u(k+1)$$

$$y(k) = 0.3y(k-1) + u(k)$$

Example 2:

$$G(z) = \frac{0.1z}{z^2 - 0.5z + 0.06}$$

$$Y(z) = \frac{0.1z}{z^2 - 0.5z + 0.06} U(z)$$



$$Y(z)(z^2 - 0.5z + 0.06) = 0.1zU(z)$$

$$y(k) = 0.5y(k-1) + 0.06y(k-2) + 0.1u(k-1)$$



# Transfer Functions In Series

$$w(k+1) = (0.43)w(k) + (0.47)u(k) \quad y(k+1) = 0.8y(k) + 0.72w(k) - 0.66w(k-1)$$



$$\frac{Y(z)}{U(z)} = \frac{W(z)}{U(z)} \frac{Y(z)}{W(z)} = G(z)H(z)$$

T.F. provide an easy way to analyze the behavior of complex structures.

# Poles of a Transfer Function

**Poles: Values of  $z$  for which the denominator is 0.**

Example:

$$H(z) = \frac{0.1z}{z^2 - 0.5z + 0.06} = \frac{z}{z - 0.3} - \frac{z}{z - 0.2}$$

Poles: 0.3, 0.2


Poles

- Determine stability
- Major effect on settling time, overshoot

$$G(z) = \frac{z}{z - a} = 1 + az^{-1} + a^2z^{-2} + a^3z^{-3} + \dots$$

- Larger  $|a|$ 
  - Slower convergence
- $|a| > 1$ 
  - Does not converge
- $a < 0$ 
  - Oscillates


# Almost All You'll Ever Need to Know About Poles



$$G(z) = \frac{z}{z - a} = 1 + az^{-1} + a^2 z^{-2} + a^3 z^{-3} + \dots$$

# Steady State Gain (ssg) of a Transfer Function

Definition and result: Steady state divided by steady state input.





$$\text{ssg of } G(z) \text{ is } \frac{y(\infty)}{u(\infty)} = G(1)$$

from final  
value thm

## Example:

$$u(\infty) = 1 \quad G(z) = \frac{z}{z - 0.5} \quad y(\infty) = 2$$




$$\frac{y(\infty)}{u(\infty)} = \frac{2}{1} = 2 = G(1)$$

# Settling Time ( $k_s$ ) of a System

Definition and result: Time until an input signal is within 2% of its steady state



from  
geometric



$$k_s \approx \frac{-4}{\log |a|}, \text{ where } |a| \text{ is the largest pole of } G(z)$$

**Examples:**


$$G(z) = \frac{z}{z-0.5} \quad k_s \approx \frac{-4}{\log 0.5} \approx 6$$




do more time steps

$$G(z) = \frac{0.34z - 0.31}{z^3 - 1.23z^2 + 0.34z},$$

poles: 0, 0.43, 0.8



$$k_s \approx \frac{-4}{\log 0.8} \approx 18$$



# Summary of Key Results

## ■ Interpretation of transfer functions

- ❖ Signal created by an impulse (impulse response)
- ❖ Sum of time delayed effects
- ❖ Encoding of a ARX model

## ■ Transfer functions in series

- ❖ End-to-end T.F. is the product of the individual T.F.


## ■ Steady state gain

- ❖ Magnitude of effect on the output for a unit change in input

## ■ Settling time

- ❖ Time until steady state is reached after applying a unit input

# Key Results for LTI Systems (LTI = Linear Time Invariant)




Stable if  $|a| < 1$ , where  $a$  is the largest pole of  $G(z)$

$k_s \approx \frac{-4}{\log |a|}$ , where  $|a|$  is the largest pole of  $G(z)$


ssg of  $G(z)$  is  $\frac{y(\infty)}{u(\infty)} = G(1)$

## Adding signals:


If  $\{a(k)\}$  and  $\{b(k)\}$  are signals, then  $\{c(k)=a(k)+b(k)\}$  has Z-Transform  $A(z)+B(z)$ .



## Transfer functions in series



is equivalent to




# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system



# Basic Controllers



## Proportional (P) Control

$$u(k) = K_p e(k)$$

$$K(z) = \frac{U(z)}{E(z)} = K_p$$

## Integral (I) Control


$$u(k+1) = K_I u(k) + e(k+1)$$

$$K(z) = \frac{U(z)}{E(z)} = \frac{K_I z}{z-1}$$

$K_p$  and  $K_I$  are called **control gains**.


# Conclusions from P vs. I Comparison

$K_p=2.3$



$r(k)=200$

$K_I=0.8$



## Conclusions:

P is fast

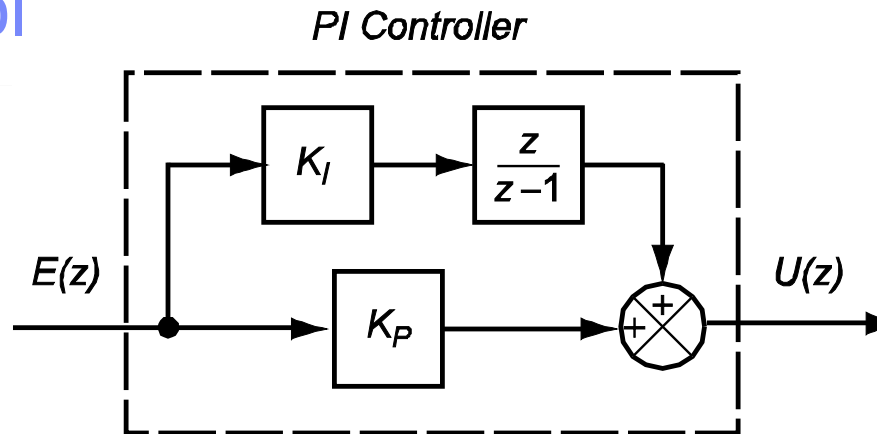
I is accurate and has less overshoot.

## Design challenge:

Make P accurate.

Reduce P's overshoot.

# PI Control



$$u(k) = u_P(k) + u_I(k)$$

$$= u(k-1) + (K_P + K_I)e(k) - K_P e(k-1)$$

$$K(z) = \frac{E(z)}{U(z)} = K_P + \frac{K_I z}{z-1} = \frac{(K_P + K_I)z - K_P}{z-1}$$

# Overview

- Open loop, closed loop
- Control system architecture
- Performance specifications (for expressing policies)
- Digital controllers
- z-Transform:
  - concept, signal composition
  - poles and stability
  - transfer function and system composition
  - proportional (P) and integral (I) control, PI control
- Effector example for a software system

# M7 Topics

## ■ DB2 Utilities Throttling (DB2 v8.1)

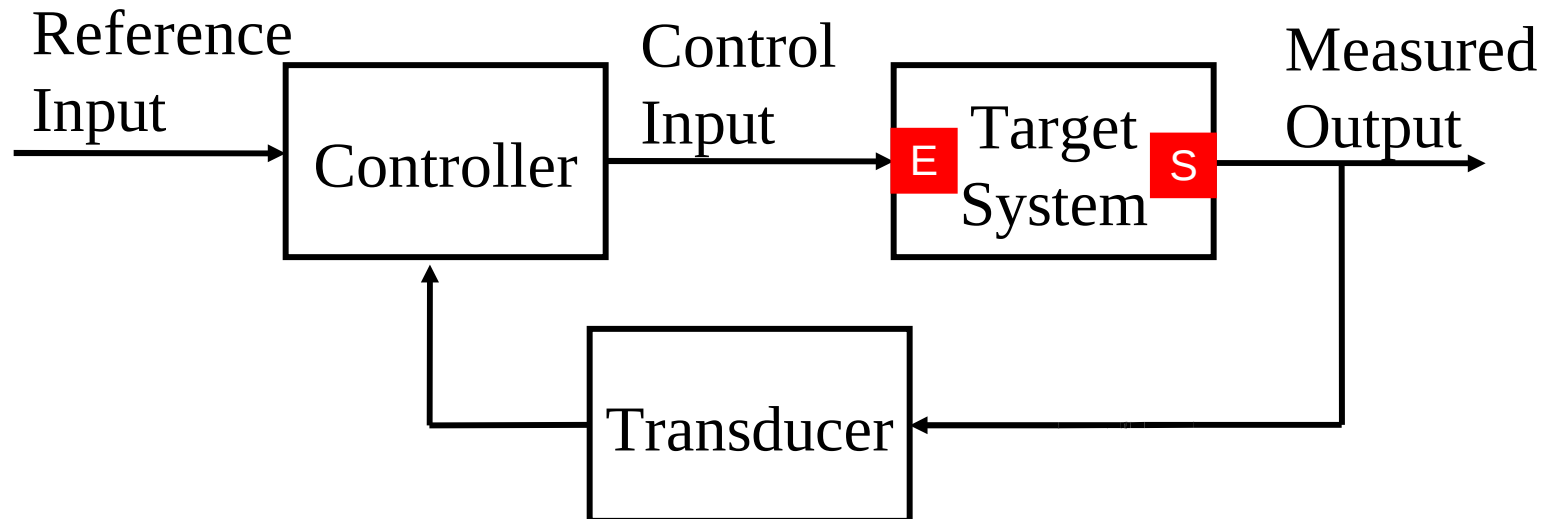
- ❖ <http://www.databasejournal.com/features/db2/article.php/3339041>
- ❖ "Throttling Utilities in the IBM DB2 Universal Database Server," Sujay Parekh, Kevin Rose, Yixin Diao, Victor Chang, Joseph L. Hellerstein, Sam Lightstone, Matthew Huras. American Control Conference, 2004.

## ■ Self-tuning memory management

- ❖ "Using MIMO Linear Control for Load Balancing in Computing Systems," Yixin Diao, Joseph L. Hellerstein, Adam Storm, Maheswaran Surendra, Sam Lightstone, Sujay Parekh, and Christian Garcia-Arellano. American Control Conference, 2004.
- ❖ "Incorporating Cost of Control Into the Design of a Load Balancing Controller," Yixin Diao, Joseph L. Hellerstein, Adam Storm, Maheswaran Surendra, Sam Lightstone, Sujay Parekh, and Christian Garcia-Arellano. Real-Time and Embedded Technology and Application Systems Symposium, 2004.

# DB2 Utilities Throttling

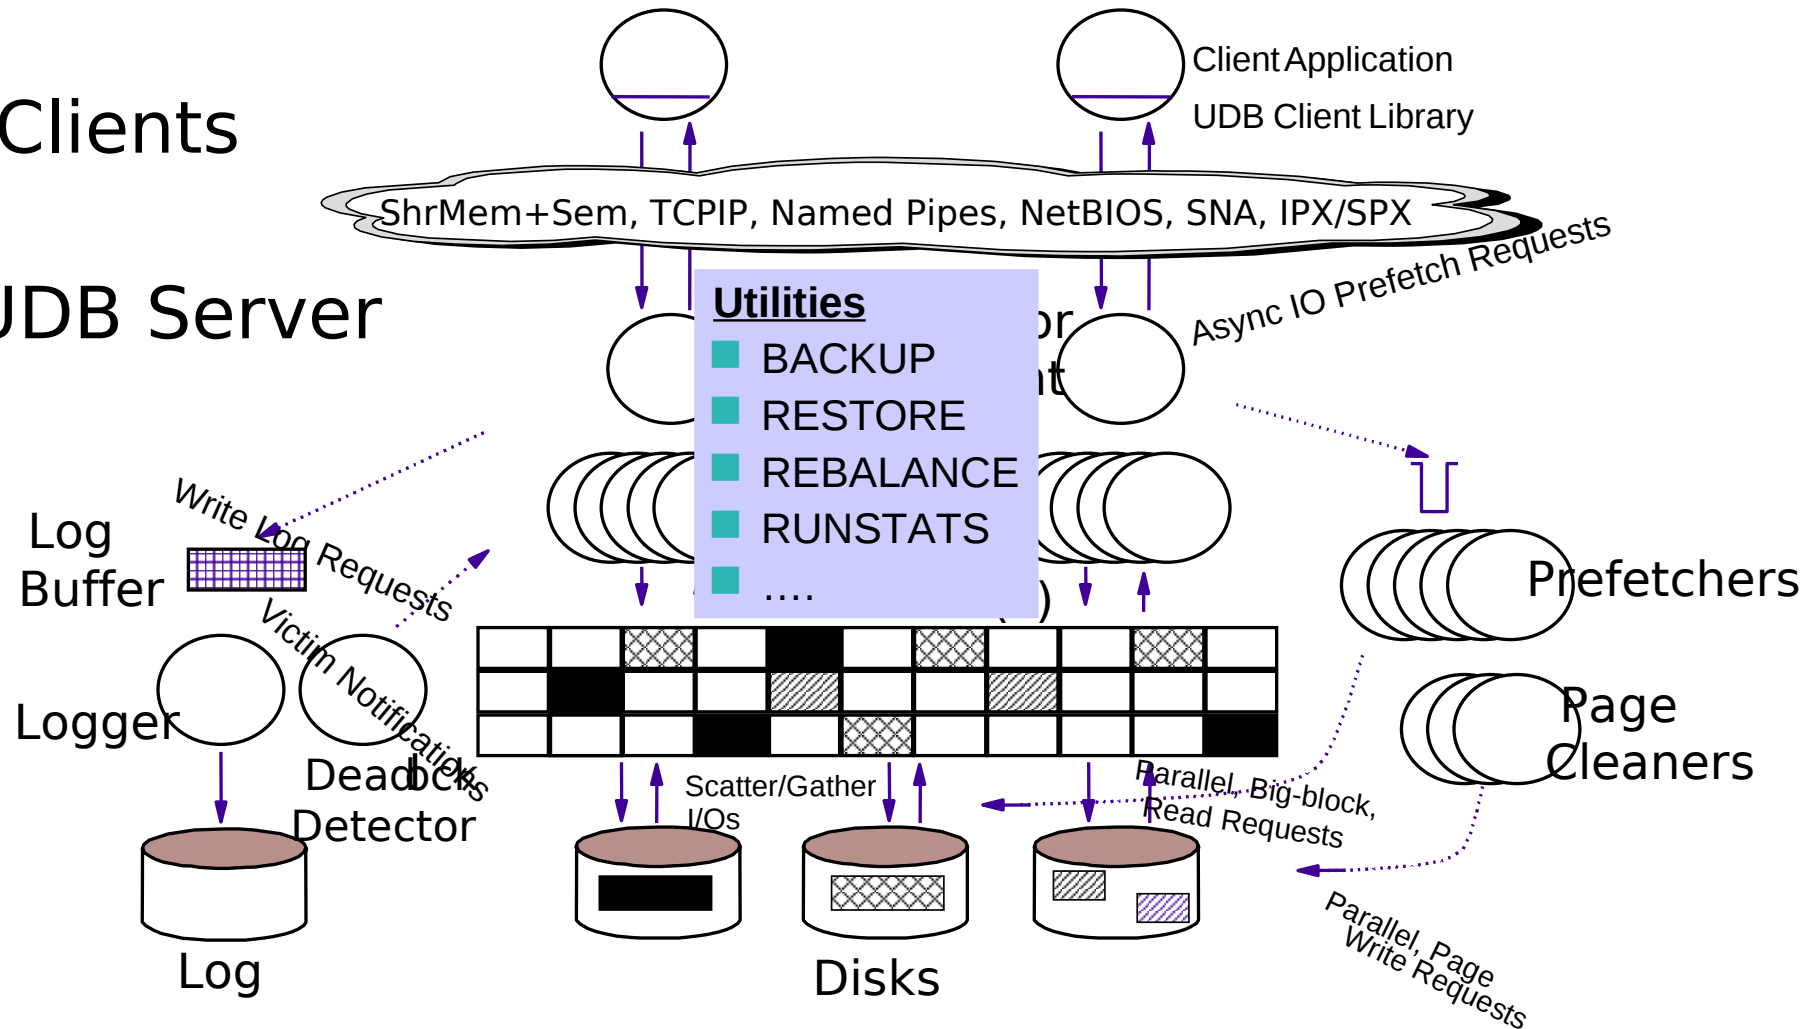
- System description
- Control problem
- Design of sensors (S) and effectors (E)
- Control system design
- Evaluation



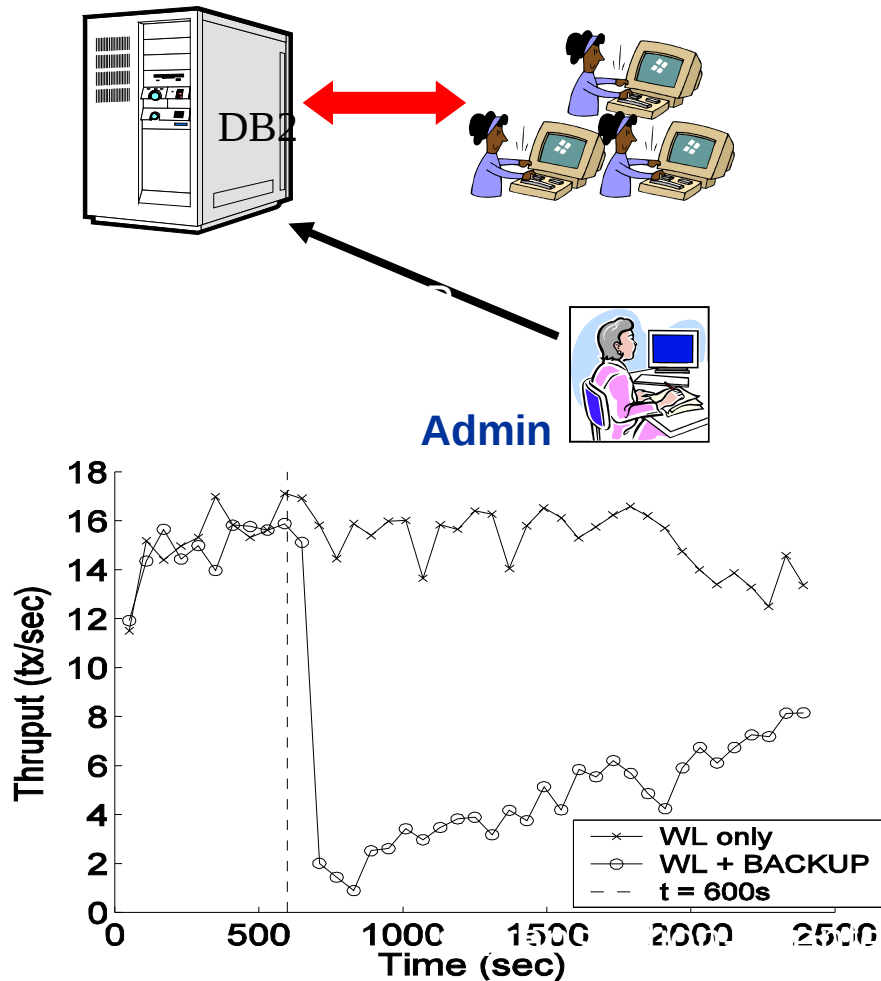
# System Description: DB/2 UDB

Clients

UDB Server




# Control Problem: Utility Impact



## Administrative policy


*There should be no more than an  $x\%$  performance degradation of production work as a result of executing administrative utilities*



**Utilities have a huge impact on production performance.**



# Effector Design: Why Not Priorities



- Assign fixed priority to Utility processes
  - ❖ Unix => higher # is lower priority
- Measure avg Workload throughput over a run
- NO EFFECT!
- Why?
  - ❖ BACKUP (and others) are I/O bound
  - ❖ OS priority == CPU priority
    - No effect on other resources
- Even if I/O priorities are present, may not solve the problem
  - ❖ Don't want to starve the utilities

# Effector Design: Self-Induced sleep (SIS)

## Augmenting the Utility

```
FUNCTION Utility()  
BEGIN  
    WHILE (NOT done)  
    BEGIN  
        ... do some work ...  
        SleepIfNeeded()  
    END  
END
```

## “Library” implementation of Throttling


```
FUNCTION SleepIfNeeded()  
BEGIN  
    (workTime, sleepTime) = GetThrottlingLevel()  
;  
    timeWorked = Now() - workStart ;  
    IF (timeWorked > workTime)  
        SLEEP( sleepTime ) ;  
        workStart = Now() ;  
    ENDIF  
END
```

**Appeal**: Handle any resource bottleneck

**Shortcoming**: Requires change to application code.



But change is modest and does not affect the customer

# Effector Design: Assessment of SIS for TPM-C



Wide range of “control volume”  
Linear effect

# Effector Design: Assessment of SIS for TPC-W



Wide range of “control volume”  
Linear effect

# Summary

- Control theory
  - a model and a tool
  - starts with system identification, „black box“ view
  - finding good effectors
  - „designing a controller“ for given policy
- Predictable outcome
  - stability analysis, controller bandwidth etc
- Is it „autonomic“?
  - adaptive, on the reactive side of the adaptivity spectrum
  - important example of the „sens/analyze/plan/exec“ cycle